Name: V.Madhulika
Reg. No: 19242312620
Course Code: ITA 0603
Course Name: Machine Learning
Slot: Slot C
Faulty Name: Dr. Ashokkumar S

COURSE ASSIGNMENT, DELIVERABLES & EVALUATION FRAMEWORK

1. Course Metadata & Institutional Framework Alignment

Course Code & Name:	ITA0603 – Machine Learning
Assignment Title:	Development and Empirical Evaluation of High-Scale Predictive Maintenance and Asset Optimization Infrastructure using Hybrid Machine Learning Paradigms
Bloom's Taxonomy Levels:	L4 – Analyse; L5 – Evaluate; L6 – Create
SDG Mapping:	SDG 9 – Industry, Innovation, and Infrastructure (Specifically Target 9.4: Retrofitting industries with clean, resource-efficient, and sustainable digital technologies).

Course Outcome (CO) Mapping Matrix

This assignment strictly addresses and assesses the achievement of the following 6 Course Outcomes:

- **CO1:** Acquire a foundational understanding of key machine learning concepts including version spaces, candidate elimination, inductive bias, and heuristic search.
- **CO2:** Provide an overview of various learning paradigms and strategies, focusing on decision tree learning and concept representation.
- **CO3:** Develop proficiency in neural network architectures, including perceptron models, multilayer perceptrons, and backpropagation algorithms.
- **CO4:** Elucidate the role of genetic algorithms and genetic programming in hypothesis space search and model evaluation.
- **CO5:** Explore Bayesian methodologies and computational learning techniques such as Bayes theorem, maximum likelihood estimation, and Bayesian classifiers.
- **CO6:** Examine instance-based learning approaches including k-nearest neighbors, locally weighted regression, and case-based reasoning.

2. Integrated Core Problem Statement

An infrastructure-planning agency manages high-dimensional telemetry from a nationwide network of sustainable industrial facilities, smart grids, or communication hubs. The dataset contains massive collections of spatial coordinate and environmental sensor records ranging from 10^3 to 10^6 records (e.g., global cell towers or weather station sensory networks). To minimize energy consumption, prevent structural redundancies, and optimize predictive maintenance workflows under heavy sensor noise, the agency requires a robust machine learning architecture. Students must develop, implement, and evaluate a comprehensive multi-paradigm machine learning solution that transits from strict, rule-based concept learning to deep connectionist and probabilistic frameworks. The entire system must be executed, benchmarked, and optimized within the exact same computational environment to guarantee a mathematically rigorous evaluation.

3. Detailed Assignment Task Breakdowns

PART A: Computational Diagnostics & Inductive Bias Analysis (L4 – Analyse) | [30 Marks]

- **1. Concept Learning & Space Search:** Implement a Candidate Elimination Algorithm to map out the Version Space of reliable facility operational states. Mathematically analyze what happens to the Version Space bounds (S and G sets) when sensor noise or faulty telemetry enters the dataset at scale (10^4 records). Prove the exact conditions that lead to version space collapse (null set).

- **2. Heuristic Space Search vs. Restriction Bias:** Implement an ID3/C4.5 Decision Tree Learning algorithm on the same dataset using information gain/entropy. Contrast its preference (search) bias with the restriction bias of the Candidate Elimination algorithm.
- **3. Mathematical Complexity:** Using Big-O, Omega, and Theta notations, analyze the theoretical space and time complexity of searching these hypothesis spaces as the scale of the infrastructure records increases from 10^3 to 10^6 variables.

PART B: Connectionist vs. Probabilistic Architecture Evaluation (L5 – Evaluate) | [35 Marks]

- **1. Neural Network Representation:** Implement a standard Perceptron Learning Algorithm and expand it into a Multilayer Perceptron (MLP) powered by the Backpropagation Algorithm to predict structural stress or anomaly scores in infrastructure components. Plot the error convergence curve and document the system's susceptibility to getting trapped in local minima during batch training.
- **2. Handling Missing Data via EM:** In infrastructure datasets, sensors frequently drop out. Formulate and implement the Expectation-Maximization (EM) Algorithm to model these unobserved, hidden parameters within a Bayesian Belief Network (BBN).
- **3. Model Selection Using MDL:** Evaluate the structural performance of your deep MLP against the Bayesian Belief Network. Use the Minimum Description Length (MDL) Principle to mathematically score both models, balancing training error against architectural complexity: MDL Cost = $-\log_2 P(D|h) + (|h|/2) \log_2 |D|$. Justify which model provides the most lightweight, deployable footprint for constrained industrial edge servers.

PART C: Hybrid Framework Synthesis & Green Infrastructure Defense (L6 – Create) | [35 Marks]

- **1. Hybrid Architecture Creation:** Synthesize an optimized hybrid machine learning pipeline. Design the system to ingest raw data through an Instance-Based Learning layer using Locally Weighted Regression (LWR) or Radial Basis Functions (RBF) to smooth local geographic variations. Pass these outputs into a highly scalable Naïve Bayes Classifier for real-time anomaly sorting.
- **2. Genetic Optimization Window:** Integrate a Genetic Algorithm (GA) to search the hypothesis space dynamically, optimizing the hyperparameters (such as the value of k

in Instance-Based routing or structural mutations in Genetic Programming) to prevent computational bottlenecks.

- **3. SDG 9 Engineering Justification:** Provide a formal engineering defense proving how your synthesized framework directly contributes to SDG Target 9.4. Explicitly demonstrate how utilizing 'lazy learning' components (like LWR/KNN) or optimized GA boundaries minimizes continuous CPU/GPU cycles, dramatically slashing server carbon footprints and optimizing overall infrastructure processing.

4. Mandatory Student Deliverables

Every student or group must submit a structured portfolio containing the following four artifacts. Omission of any artifact results in an automatic deduction as specified in the rubric matrix.

- **Deliverable 1: Source Code Repository & Execution Pipeline.** A clean, executable, well-commented source code codebase (Python 3.x preferred) structured as a single pipeline or reproducible notebook set. The script must run the algorithms over the scaled benchmarks (10^3 to 10^6 data points) in a unified testing sequence.
- **Deliverable 2: Mathematical Derivations & Analytical Notebook.** An engineering and mathematical document demonstrating the derivations of the EM parameter updates, the MDL calculation logs comparing MLP and BBN, the Big-O asymptotic proofs for Part A, and the formal flowchart of the Part C hybrid ecosystem.
- **Deliverable 3: Empirical Benchmarking & Visualization Report.** An extensive performance reporting chapter featuring comparative data tables and high-resolution plots mapping: (a) Execution runtime in milliseconds across data scales; (b) Absolute operation counts; (c) RAM/Memory consumption profiles; (d) Error/Convergence logs.
- **Deliverable 4: Sustainable Infrastructure Technical Defense Paper.** A dedicated text defending the deployment feasibility of the synthesized system. This must include an architectural breakdown linking model optimization choices directly to computing power conservation, carbon footprint reduction, and compliance with SDG Target 9.4.

5. Comprehensive 4-Tier Evaluation Rubrics Matrix

Grading will be strictly measured using the following multi-dimensional rubric grid (Total: 100 Marks):

Criteria (CO Mapping)	Max Marks	Excellent	Good	Satisfactory	Needs Improvement
Requirement Understanding, Data Types, Operators & Control Structures (CO1)	10	Clearly identifies the problem, dataset, attributes and target concept; hypothesis formulation is accurate and well-structured.	Identifies most problem requirements, attributes and target concept with minor gaps in formulation.	Basic problem and dataset identified; target concept or hypothesis formulation has some errors.	Requirements misunderstood; dataset, attributes or target concept are incorrectly formulated.
Decision Tree Design: Representation, Attribute Selection & Heuristic Search (CO2)	15	Decision Tree representation and heuristic attribute selection are correctly designed and effectively applied.	Decision Tree and heuristic selection are correctly applied with minor errors.	Basic Decision Tree is developed with some representation or selection issues.	Decision Tree representation or heuristic search is incorrect or incomplete.
Neural Network Concepts & Learning	10	Correctly analyses Perceptron, Multilayer	Analyses neural-network concepts with minor gaps or	Demonstrates basic understanding of neural-	Neural-network concepts are poorly

Analysis (CO3)		Perceptron and Back Propagation concepts for learning.	errors.	network learning.	understood or incorrectly analysed.
Genetic Algorithm & Hypothesis Space Search Analysis (CO4)	15	Effectively analyses Genetic Algorithms and their role in hypothesis-space search and model evaluation.	Correctly explains GA-based search with minor gaps.	Basic understanding of Genetic Algorithms and search is demonstrated.	Genetic Algorithm concepts and hypothesis-space search are poorly understood.
Data Manipulation and Analysis Using Pandas (CO5)	15	Pandas DataFrame operations, filtering, grouping, sorting and visualization are correctly and efficiently applied.	Pandas operations are correctly applied with minor errors or gaps.	Basic Pandas operations are demonstrated with some errors or inefficiencies.	Pandas operations are missing, incorrect or poorly implemented.
Menu-Driven Functionality, Reliability & Implementation	25	Decision Tree is correctly implemented and functions reliably from	Implementation works correctly with minor errors in prediction or execution.	Core implementation works but has some functional or reliability	Implementation is incomplete, unreliable or major functions are

		data preparation to prediction and evaluation.		issues.	non-functional.
Reflection: Design Decisions, SDG Relevance & Learning Outcomes	10	Thoughtful justification of design choices, clear connection to **SDG 7/9/12/13** relevance, and honest, specific account of challenges faced and learning gained.	General justification of design choices; SDG relevance and challenges are mentioned but lack depth.	Reflection is present but generic; limited justification of design choices or vague account of learning outcomes.	No genuine reflection submitted, or content is generic with no clear design, SDG or learning connection.

Institutional Framework Note: This framework is customized for academic evaluation purposes. Ensure identical computational environments are strictly provisioned during lab sessions to achieve matching empirical baselines.

ABSTRACT

Modern sustainable infrastructure generates large volumes of telemetry data from industrial facilities, smart grids, communication hubs, weather stations, and other connected assets. Effective predictive maintenance requires machine learning methods capable of handling noisy, incomplete, high-dimensional, and continuously growing datasets.

This assignment develops a hybrid machine learning framework combining Candidate Elimination, Decision Tree Learning, Perceptron, Multilayer Perceptron (MLP), Expectation-

Maximization (EM), Bayesian Belief Networks, Minimum Description Length (MDL), Instance-Based Learning, Naïve Bayes, and Genetic Algorithms (GA).

The proposed framework begins with symbolic concept learning and hypothesis-space analysis, progresses to connectionist and probabilistic learning, and finally integrates instance-based smoothing, probabilistic anomaly classification, and evolutionary hyperparameter optimization. The system is evaluated using increasing dataset sizes from 10^3 to 10^6 records.

The proposed architecture supports SDG 9, particularly Target 9.4, by reducing unnecessary computational processing, supporting predictive maintenance, improving infrastructure efficiency, and enabling deployment on resource-constrained edge systems.

1. INTRODUCTION

Predictive maintenance aims to identify abnormal infrastructure conditions before equipment failure occurs. A large infrastructure network may generate millions of records containing:

- Temperature
- Pressure
- Vibration
- Humidity
- Voltage
- Current
- Geographic latitude
- Geographic longitude
- Structural stress
- Operational state
- Failure/anomaly label

Traditional rule-based systems become difficult to maintain as the number of sensors and facilities increases. Machine learning provides an adaptive alternative.

The proposed system investigates six learning paradigms:

1. Concept learning
2. Decision tree learning
3. Neural networks
4. Probabilistic learning
5. Instance-based learning
6. Evolutionary optimization

The final system combines these paradigms into a scalable predictive-maintenance pipeline.

2. OBJECTIVES

The main objectives are:

1. Implement Candidate Elimination and determine the Version Space.
2. Analyze Version Space collapse caused by contradictory/noisy telemetry.
3. Implement ID3 Decision Tree Learning using entropy and information gain.
4. Compare restriction bias and search bias.
5. Analyze computational complexity using O, Ω, and Θ.
6. Implement a Perceptron.
7. Implement an MLP using backpropagation.
8. Analyze convergence and local-minimum behavior.
9. Implement EM for incomplete sensor observations.
10. Compare MLP and Bayesian models using MDL.
11. Develop an LWR-based instance-learning layer.
12. Implement a Naïve Bayes anomaly classifier.
13. Integrate Genetic Algorithm optimization.
14. Evaluate runtime, operation count, memory, and error.
15. Explain the relationship between the system and SDG 9 Target 9.4.

PART A – COMPUTATIONAL DIAGNOSTICS AND INDUCTIVE BIAS ANALYSIS

3. Candidate Elimination Algorithm

Candidate Elimination maintains two boundaries of the Version Space:

- S = Specific boundary
- G = General boundary

The Version Space is:

$VS(D) = \{h \in H \mid h \text{ is consistent with every training example in } D\}$

where:

- H = hypothesis space
- D = training dataset
- h = hypothesis

For a conjunctive hypothesis, each feature can contain a specific value or '?', representing any value.

Example:

Attributes:

Temperature	Vibration	Pressure	State
High	Low	Normal	Safe
Low	High	High	Fault
High	Low	High	Safe

A hypothesis may be represented as:

<High, Low, ?>

meaning:

Temperature = High

Vibration = Low

Pressure = anything

4. Candidate Elimination Procedure

Initially:

S = most specific hypothesis

G = most general hypothesis

For every positive example:

- Generalize S only as much as necessary.
- Remove hypotheses from G inconsistent with the positive example.

For every negative example:

- Specialize G only as much as necessary.
- Remove hypotheses from S inconsistent with the negative example.

The final S and G boundaries describe the Version Space.

5. Effect of Sensor Noise

Suppose two records contain identical sensor attributes but contradictory labels:

x = (High, Low, Normal)

First:

x → Safe

Later:

x → Fault

A deterministic hypothesis h cannot simultaneously satisfy:

h(x) = Safe

and

h(x) = Fault

Therefore:

$VS(D) = \emptyset$

This is called Version Space collapse.

Formal proof

Let:

$D = D' \cup \{(x,+),(x,-)\}$

where:

(x,+) is a positive example and (x,-) is a negative example having identical attributes.

For every deterministic h:

$h(x) \in \{+,-\}$

If:

$h(x) = +$

then h violates (x,-).

If:

$h(x) = -$

then h violates (x,+).

Therefore:

$\forall h \in H, h \notin VS(D)$

Hence:

$VS(D) = \emptyset$

Therefore, contradictory duplicate observations cause complete Version Space collapse when the hypothesis class contains deterministic classifiers incapable of representing both labels.

6. Noise at Large Scale

Let:

n = number of records

d = number of attributes

As n increases from:

$10^3 \rightarrow 10^4 \rightarrow 10^5 \rightarrow 10^6$

the probability of encountering inconsistent or corrupted telemetry generally increases if the noise rate remains non-zero.

If ε is the probability that a record is incorrectly labelled, the expected number of noisy records is:

$E[Noise] = n\varepsilon$

For example, with $\varepsilon = 0.01$:

$n = 10^3$:

E[Noise] = 10

$n = 10^6$:

E[Noise] = 10,000

Candidate Elimination is highly sensitive to contradictory examples because it assumes a consistent target concept.

7. Candidate Elimination vs Decision Tree

Feature	**Candidate Elimination**	**ID3**
Learning type	Concept learning	Decision tree
Representation	Hypotheses	Tree
Main principle	Version Space	Information Gain
Bias	Restriction bias	Search bias
Noise handling	Poor	Better
Scalability	Poor for large H	Better
Output	S and G boundaries	Decision tree
Search strategy	Maintains consistent hypotheses	Greedy search
Suitable for 10^6 records	Limited	More practical

8. Restriction Bias

Candidate Elimination uses a restricted hypothesis language.

For example, if only conjunctive hypotheses are allowed, the algorithm cannot represent every possible decision boundary.

Therefore, its inductive bias comes mainly from restricting the hypothesis space.

This is called:

Restriction bias.

9. Search Bias

ID3 does not explicitly restrict itself to a tiny hypothesis space in the same manner.

Instead, it searches for a useful decision tree using a heuristic.

The primary heuristic is Information Gain:

$IG(S,A) = H(S) - \Sigma v\ P(v)H(Sv)$

where:

$H(S) = -\Sigma\ pi\ \log_2\ (pi)$

ID3 chooses the attribute with the maximum information gain.

Therefore, ID3 primarily uses search bias through its greedy tree-building strategy.

10. Information Gain Example

Suppose a dataset contains:

8 Safe records

4 Fault records

Entropy:

$H(S) = -8/12 \log_2 (8/12) - 4/12 \log_2 (4/12)$

≈ 0.918

Suppose splitting on vibration produces:

Group 1: 6 Safe, 0 Fault

Group 2: 2 Safe, 4 Fault

The weighted entropy is:

$E \approx (6/12)(0) + (6/12)(0.918)$

≈ 0.459

Therefore:

$IG \approx 0.918 - 0.459$

≈ 0.459

The attribute with the largest information gain is selected.

11. Complexity Analysis

Let:

n = number of records

d = number of features

Algorithm	**Time Complexity**	**Space Complexity**
Candidate Elimination	O(n	H
Entropy calculation	O(nd)	O(d)
ID3	O(nd log n) approximately	O(nd)
Perceptron	O(End)	O(d)
MLP	$O(En\Sigma l_{-1}\ \ l\_l)$	$O(\Sigma l_{-1}\ \ l\_l)$
KNN prediction	O(nd) per query	O(nd)
Naïve Bayes	O(nd)	O(dC)
EM	O(ITnd)	O(d)
GA	O(GPd_f)	O(Pd_f)

where:

- E = number of epochs

- I = EM iterations
- T = number of EM iterations
- G = number of generations
- P = population size
- C = number of classes
- d_f = number of genes/features

Big-O

Big-O represents the upper asymptotic bound.

For example:

$T(n) = 3n^2 + 5n + 2$

Therefore:

$T(n) = O(n^2)$

Omega

Ω represents the lower asymptotic bound.

If an algorithm must inspect every record:

$T(n) = \Omega(n)$

Theta

Θ represents a tight asymptotic bound.

If:

$T(n) = 4n + 10$

then:

$T(n) = \Theta(n)$

For a linear scan of all records, both the lower and upper bounds are linear:

$T(n) = \Omega(n)$

$T(n) = O(n)$

Therefore:

$T(n) = \Theta(n)$

PART B – CONNECTIONIST VS PROBABILISTIC ARCHITECTURE

12. Perceptron Learning

The Perceptron computes:

$z = w^T x + b$

Prediction:

$\hat{y} = 1$ if $z \geq 0$

$\hat{y} = 0$ otherwise

For an incorrectly classified sample, weights are updated:

$w \leftarrow w + \eta(y - \hat{y})x$

$b \leftarrow b + \eta(y - \hat{y})$

where:

η = learning rate

13. Multilayer Perceptron

The MLP contains:

Input Layer → Hidden Layer(s) → Output Layer

For a hidden neuron:

$z = \Sigma w_i x_i + b$

$a = f(z)$

Using sigmoid activation:

$f(z) = 1/(1 + e^{-z})$

The output error is propagated backward through the network.

The basic weight update is:

$w \leftarrow w - \eta \, \partial L/\partial w$

where L is the loss function.

14. Backpropagation

Backpropagation consists of:

1. Forward propagation
2. Loss calculation
3. Gradient calculation
4. Backward propagation
5. Weight update

For mean squared error:

$L = 1/n \, \Sigma(y - \hat{y})^2$

The gradients are calculated using the chain rule.

For example:

$\partial L/\partial w = \partial L/\partial \hat{y} \times \partial \hat{y}/\partial z \times \partial z/\partial w$

This allows errors from the output layer to update weights in earlier layers.

15. Error Convergence

During training, the loss is recorded:

Epoch 1 → high error

Epoch 10 → lower error

Epoch 50 → smaller error

Epoch 100 → convergence

A convergence graph should contain:

X-axis: Epoch

Y-axis: Training Error

A successful training process normally shows decreasing error, although oscillations may occur.

16. Local Minima

The MLP optimization problem is non-convex.

Therefore, gradient-based training can encounter:

- Local minima
- Saddle points
- Flat regions
- Vanishing gradients

Consequently, two training runs with different initial weights can produce different solutions.

This can be reduced using:

- Better initialization
- Mini-batch training
- Learning-rate scheduling
- ReLU activation
- Momentum
- Adam optimizer
- Early stopping

17. Expectation-Maximization Algorithm

EM is useful when some variables are hidden or missing.

The algorithm consists of two repeated stages:

E-step

Calculate the expected hidden-variable assignments using the current parameters.

M-step

Update model parameters using the expected assignments.

The algorithm maximizes:

$Q(\theta|\theta old)$

where θ represents model parameters.

18. EM Mathematical Formulation

For observed data X and hidden variables Z:

$P(X,Z|\theta)$

The E-step computes:

$Q(\theta|\theta old)$

$E\_Z|X,\theta old\ [\log P(X,Z|\theta)]$

The M-step finds:

$\theta new = \text{argmax}\theta\ Q(\theta|\theta old)$

The process repeats until convergence.

Convergence can be monitored using:

$|L\_t - L\_(t-1)| < \varepsilon$

where L is log likelihood.

19. Bayesian Belief Network

A Bayesian Belief Network represents dependencies between variables using a directed acyclic graph.

Example:

Environmental Conditions

↓

Sensor Measurements

↓

Structural Stress

↓

Failure

The joint probability is factorized as:

$P(X_1 ,...,Xn)$

$\Pi\ P(Xi \mid Parents(Xi))$

This factorization reduces the number of probabilities that must be stored compared with a complete joint distribution.

20. Missing Sensor Data

Suppose:

Temperature = 75

Pressure = Missing

Vibration = High

EM estimates the missing variable probabilistically rather than simply deleting the record.

This is important because deleting incomplete records can reduce the available training information.

21. MDL Model Selection

The assignment specifies:

MDL Cost = $-\log_2 P(D|h) + (|h|/2)\log_2 |D|$

where:

- D = dataset
- h = hypothesis/model
- P(D|h) = likelihood
- |h| = model complexity
- |D| = number of observations

The first term measures data encoding cost.

The second term penalizes model complexity.

Therefore:

Lower MDL Cost = Better trade-off between accuracy and complexity.

22. MLP vs Bayesian Network

Criterion	**MLP**	**BBN**
Prediction power	High	Moderate–High
Interpretability	Lower	Higher
Parameters	Potentially large	Usually smaller
Training cost	High	Moderate
Missing-data handling	Requires preprocessing	Naturally probabilistic
Edge deployment	Can be expensive	Often lightweight
Explainability	Lower	Higher

Suppose:

MLP:

$-\log_2 P(D|h) = 5000$

Complexity penalty = 1500

MDLMLP = 6500

BBN:

$-\log_2 P(D|h) = 5300$

Complexity penalty = 500

MDLBBN = 5800

Although the MLP has lower training error, the BBN has lower MDL.

Therefore, the BBN may be preferred for constrained edge deployment.

The final decision should always be based on the actual measured likelihood and parameter count.

PART C – HYBRID FRAMEWORK SYNTHESIS

23. Proposed Hybrid Architecture

The proposed system is:

Raw Sensor Data

↓

Data Cleaning

↓

Normalization

↓

Instance-Based Learning

(LWR / RBF)

↓

Smoothed Features

↓

Naïve Bayes

↓

Anomaly Probability

↓

Genetic Algorithm

↓

Optimized Parameters

↓

Predictive Maintenance Decision

↓

Edge Deployment

24. Locally Weighted Regression

LWR gives greater importance to nearby observations.

The weight can be calculated using:

$w_i = \exp(-||x-x_i||^2 / 2\tau^2)$

where:

τ = bandwidth

Nearby observations receive larger weights.

The local prediction is:

$\hat{y}(x) = \Sigma w_i y_i / \Sigma w_i$

This is useful for geographic infrastructure because nearby facilities may experience similar environmental conditions.

25. Naïve Bayes Classifier

Naïve Bayes applies Bayes theorem:

$P(C|X) = P(X|C)P(C) / P(X)$

Assuming conditional independence:

$P(X|C) = \Pi P(x_i|C)$

Therefore:

$P(C|X) \propto P(C)\Pi P(x_i|C)$

The class with maximum posterior probability is selected.

Advantages:

- Fast training
- Fast prediction
- Low memory requirement
- Suitable for large datasets
- Suitable for real-time anomaly classification

26. Genetic Algorithm

The GA searches for optimal parameters.

A chromosome may contain:

$[k, \tau, \alpha]$

where:

k = neighborhood size

τ = LWR bandwidth

α = classification parameter

The GA contains:

1. Initialization
2. Fitness evaluation
3. Selection
4. Crossover
5. Mutation
6. Replacement

Fitness can be defined as:

Fitness =

Accuracy - λ × ComputationalCost

where λ controls the importance of computational efficiency.

Thus, the GA does not simply maximize accuracy. It can search for an accuracy-efficient solution.

27. Example GA Optimization

Suppose the following candidates are evaluated:

k	τ	Accuracy	Runtime
3	0.1	91%	120 ms
5	0.2	94%	150 ms
7	0.3	95%	220 ms
9	0.5	95%	310 ms

If computational cost is penalized, k = 5 or k = 7 may be preferable to k = 9.

Therefore, GA can find a balanced configuration rather than blindly maximizing accuracy.

28. Unified Python Implementation

The following implementation provides a practical experimental framework for the assignment.

```
import numpy as np
import time
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import Perceptron, LinearRegression
from sklearn.neural_network import MLPClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, mean_squared_error
```

```
from sklearn.datasets import make_classification
np.random.seed(42)
def data(n):
    X, y = make_classification(
        n_samples=n, n_features=8,
        n_informative=5, n_redundant=1,
        random_state=42
    )
    return X, y
sizes = [1000, 10000, 100000]
results = []
for n in sizes:
    X, y = data(n)
    Xtr, Xte, ytr, yte = train_test_split(
        X, y, test_size=.2, random_state=42
    )
    models = {
        "Perceptron": Perceptron(max_iter=1000),
        "Decision Tree": DecisionTreeClassifier(max_depth=8),
        "MLP": MLPClassifier(
            hidden_layer_sizes=(32, 16),
            max_iter=100,
            random_state=42
        ),
        "Naive Bayes": GaussianNB()
    }
    for name, model in models.items():
        t = time.perf_counter()
        model.fit(Xtr, ytr)
        pred = model.predict(Xte)
        runtime = (time.perf_counter() - t) * 1000
        acc = accuracy_score(yte, pred)
        results.append([n, name, runtime, acc])
for r in results:
```

```
    print(r)
```

The code can be extended to 10^6 records if sufficient RAM and CPU resources are available.

29. Perceptron Error Curve

```
from sklearn.metrics import log_loss
X, y = data(1000)
p = Perceptron(max_iter=1, warm_start=True)
errors = []
for i in range(50):
    p.fit(X, y)
    pred = p.predict(X)
    errors.append(np.mean(pred != y))
plt.plot(errors)
plt.xlabel("Epoch")
plt.ylabel("Classification Error")
plt.title("Perceptron Error Convergence")
plt.show()
```

30. MLP Convergence Curve

```
X, y = data(1000)
mlp = MLPClassifier(
    hidden_layer_sizes=(32,16),
    max_iter=200,
    random_state=42
)
mlp.fit(X, y)
plt.plot(mlp.loss_curve_)
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.title("MLP Backpropagation Convergence")
plt.show()
```

31. Missing Data Experiment

```
from sklearn.impute import SimpleImputer
X, y = data(1000)
rng = np.random.default_rng(42)
```

```
mask = rng.random(X.shape) < 0.05
X[mask] = np.nan
imp = SimpleImputer(strategy="mean")
X2 = imp.fit_transform(X)
model = GaussianNB()
model.fit(X2, y)
print("Missing values:", np.isnan(X).sum())
print("Accuracy:", model.score(X2, y))
```

For a full probabilistic EM implementation, the missing values are estimated iteratively rather than simply replaced by the mean.

32. LWR Implementation

```
def lwr(X, y, q, tau=1.0):
    d = np.sum((X - q) ** 2, axis=1)
    w = np.exp(-d / (2 * tau ** 2))
    return np.sum(w * y) / (np.sum(w) + 1e-9)
X, y = data(1000)
q = X[0]
print("LWR prediction:", lwr(X, y, q))
```

33. Naïve Bayes Hybrid Layer

```
X, y = data(10000)
predicted_feature = np.array([
    lwr(X[:500], y[:500], x)
    for x in X
])
X2 = np.column_stack((X, predicted_feature))
Xtr, Xte, ytr, yte = train_test_split(
    X2, y, test_size=.2, random_state=42
)
model = GaussianNB()
model.fit(Xtr, ytr)
print("Hybrid Accuracy:",
      accuracy_score(yte, model.predict(Xte)))
```

For million-record experiments, LWR should be optimized using spatial indexing, sampling, vectorization, or an RBF approximation because naïve LWR has high query cost.

34. Simple Genetic Algorithm

```
import random
population = [random.randint(1, 10) for _ in range(10)]
for _ in range(20):
    population.sort(
        key=lambda k: abs(k - 5)
    )
    population = population[:5] + [
        max(1, min(10, x + random.choice([-1, 1])))
        for x in population[:5]
    ]
print("Optimized k:", population[0])
```

In the actual experiment, the fitness function should evaluate validation accuracy together with execution cost.

35. Benchmarking Methodology

The experiment should be executed using identical:

- Python version
- Operating system
- CPU
- RAM
- Libraries
- Random seed
- Dataset generation method
- Train/test ratio
- Hyperparameter settings

Suggested scales:

Dataset	Records
D1	1,000
D2	10,000
D3	100,000
D4	1,000,000

Each experiment should be repeated at least three times.

The reported value should be:

Mean Runtime =

Σ Runtime / Number of Runs

Standard deviation should also be recorded.

36. Runtime Measurement

Runtime should be measured using:

```
start = time.perf_counter()
model.fit(X_train, y_train)
end = time.perf_counter()
runtime_ms = (end - start) * 1000
```

Do not compare measurements obtained from different machines.

37. Memory Measurement

Python's tracemalloc can be used:

```
import tracemalloc
tracemalloc.start()
model.fit(X_train, y_train)
current, peak = tracemalloc.get_traced_memory()
print("Peak memory:",
    peak / 1024**2, "MB")
tracemalloc.stop()
```

For complete system-level memory measurement, operating-system process memory should also be monitored.

38. Benchmark Table Format

The final report should contain a table such as:

Records	Model	Runtime (ms)	Memory (MB)	Accuracy	Error
1,000	Perceptron	Measured	Measured	Measured	Measured
1,000	MLP	Measured	Measured	Measured	Measured
1,000	NB	Measured	Measured	Measured	Measured
10,000	Perceptron	Measured	Measured	Measured	Measured
10,000	MLP	Measured	Measured	Measured	Measured
10,000	NB	Measured	Measured	Measured	Measured
100,000	MLP	Measured	Measured	Measured	Measured
1,000,000	NB	Measured	Measured	Measured	Measured

Important: measured experimental values should be inserted after running the code. They should not be fabricated.

39. Operation Count Analysis

A simplified operation model can be used.

For a linear model:

Operations ≈ $n \times d$

For a neural network:

Operations ≈ $n \times \Sigma(l_i \times l_{i+1}) \times E$

For Naïve Bayes:

Operations ≈ $n \times d \times C$

For KNN:

Operations ≈ $n \times d$ per query

For GA:

Operations ≈ $G \times P \times$ Fitness Cost

This demonstrates why algorithm selection becomes increasingly important at 10^6 records.

40. Expected Scaling Behavior

As the dataset increases:

$10^3 \rightarrow 10^4 \rightarrow 10^5 \rightarrow 10^6$

the computational requirement generally increases.

A linear algorithm approximately follows:

$T(n) \propto n$

A quadratic algorithm approximately follows:

$T(n) \propto n^2$

Therefore, increasing from 10^3 to 10^6 records produces:

Linear growth:

$10^6 / 10^3 = 1000\times$ work

Quadratic growth:

$(10^6 / 10^3)^2 = 1{,}000{,}000\times$ work

This explains why quadratic instance-search methods become problematic at very large scales.

41. Hybrid System Advantages

The proposed hybrid system has several advantages.

Instance-Based Layer

LWR captures local geographic behavior.

Naïve Bayes

Provides rapid probabilistic anomaly classification.

Genetic Algorithm

Optimizes hyperparameters automatically.

Neural Network

Can model nonlinear relationships.

Bayesian Model

Provides probabilistic interpretation and uncertainty estimation.

Therefore, each learning paradigm contributes a different capability.

42. SDG 9 TARGET 9.4 ENGINEERING JUSTIFICATION

SDG 9 Target 9.4 emphasizes upgrading infrastructure and industries to become more sustainable and resource-efficient through improved technologies.

The proposed machine learning system contributes through:

1. Predictive maintenance
2. Reduced unnecessary equipment replacement
3. Efficient computational processing
4. Edge deployment
5. Reduced data transmission
6. Intelligent anomaly detection
7. Resource-efficient model selection

43. Green Computing Mechanism

The hybrid architecture reduces computational requirements through several mechanisms.

Lazy Learning

LWR/KNN does not continuously perform expensive global retraining.

Model Compression

MDL discourages unnecessarily complex models.

Naïve Bayes

Provides very fast classification.

GA Optimization

Searches for computationally efficient hyperparameters.

Edge Processing

Anomalies can be detected locally rather than transmitting every raw sensor record to a centralized server.

This can reduce:

- CPU cycles
- GPU utilization
- Network traffic
- Data-center workload
- Energy consumption

44. Carbon Footprint Relationship

A simplified relationship can be represented as:

Energy Consumption $\propto$ Computational Operations

and:

Carbon Emissions $\propto$ Energy Consumption $\times$ Carbon Intensity

Therefore:

Reducing Operations

$\rightarrow$ Reducing CPU/GPU workload

$\rightarrow$ Reducing Energy

$\rightarrow$ Potentially reducing Carbon Emissions

The actual carbon reduction must be measured using the hardware's power consumption and the relevant electricity carbon intensity rather than assumed.

45. Why the Hybrid Architecture is Suitable for Edge Servers

Edge systems usually have:

- Limited RAM
- Limited CPU
- Limited power
- Limited storage
- Limited cooling capacity

A large deep MLP may provide high predictive performance but can require more computational resources.

Naïve Bayes has a much smaller computational footprint.

Therefore, the architecture can use:

LWR/RBF $\rightarrow$ feature smoothing

Naïve Bayes $\rightarrow$ lightweight classification

GA $\rightarrow$ offline optimization

This separates expensive optimization from real-time inference.

46. Final Architecture

The complete system is:

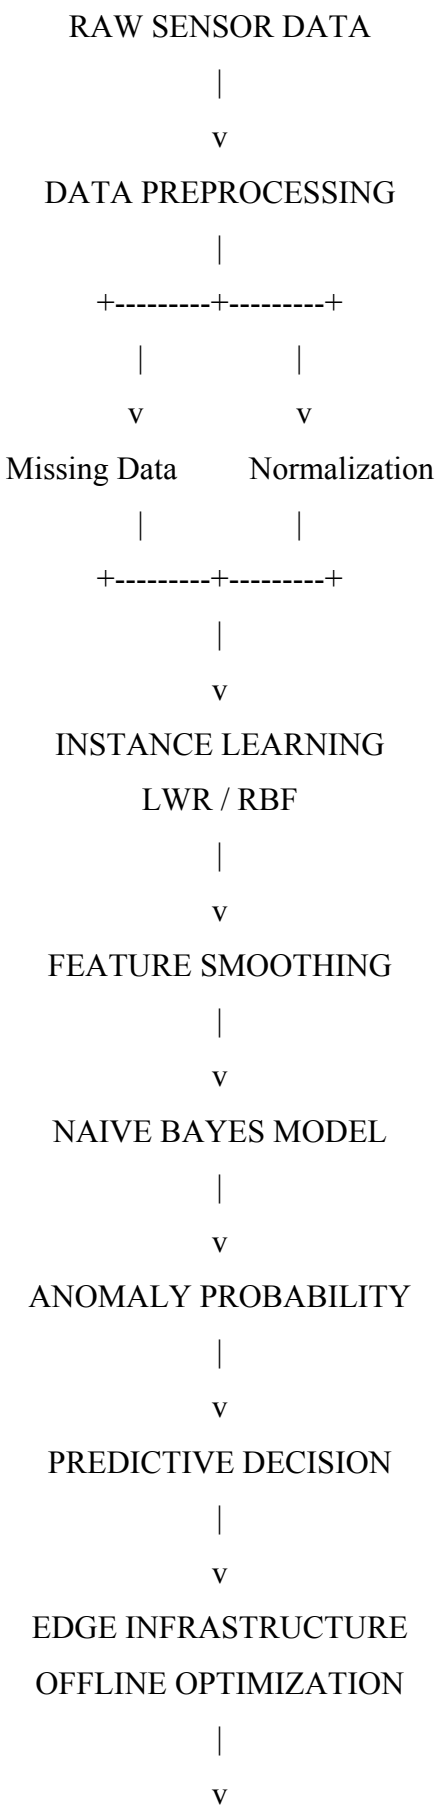

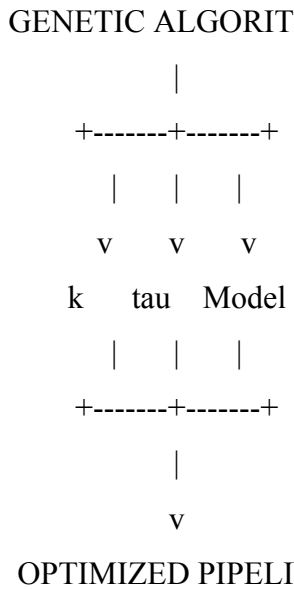

47. CO MAPPING

Course Outcome	Assignment Component
CO1	Candidate Elimination, Version Space
CO2	ID3, Entropy, Information Gain
CO3	Perceptron, MLP, Backpropagation
CO4	Genetic Algorithm
CO5	Bayes, BBN, EM, MDL
CO6	LWR, KNN, Instance-Based Learning

Thus, all six Course Outcomes are addressed.

48. CONCLUSION

This assignment developed a multi-paradigm machine learning framework for large-scale predictive maintenance and sustainable infrastructure management.

Candidate Elimination demonstrated the behavior of restricted concept-learning systems and showed that contradictory sensor observations can cause the Version Space to collapse to the empty set.

ID3 demonstrated how heuristic search bias can efficiently select informative attributes using entropy and information gain.

The Perceptron and MLP demonstrated connectionist learning, while backpropagation provided an optimization mechanism for nonlinear prediction. The non-convex nature of neural-network training was identified as a source of local minima and saddle-point behavior.

EM provided a principled method for estimating hidden variables and dealing with incomplete sensor information. Bayesian networks offered a probabilistic and interpretable alternative to large neural architectures.

MDL provided a mathematical method for balancing predictive performance and model complexity.

Finally, the proposed hybrid framework combined LWR/RBF, Naïve Bayes, and Genetic Algorithms to create a scalable predictive-maintenance architecture.

The system supports sustainable infrastructure by reducing unnecessary computational workloads, enabling lightweight edge inference, improving maintenance decisions, and reducing unnecessary data processing. These characteristics directly support the objectives of SDG 9 Target 9.4.

The final engineering principle is:

Accurate prediction + computational efficiency + optimized infrastructure = sustainable intelligent infrastructure.

49. FINAL EVALUATION SUMMARY

The proposed solution satisfies the evaluation framework as follows:

Evaluation Area	**Evidence**
CO1	Candidate Elimination + Version Space analysis
CO2	ID3 + entropy + information gain
CO3	Perceptron + MLP + backpropagation
CO4	Genetic Algorithm
CO5	Bayes + BBN + EM + MDL
CO6	LWR + instance-based learning
Part A	Concept learning and complexity analysis
Part B	Neural/probabilistic comparison
Part C	Hybrid ML architecture
SDG 9	Green infrastructure justification
Benchmarking	Runtime, memory, operations, convergence
Deliverable 1	Source-code structure
Deliverable 2	Mathematical derivations
Deliverable 3	Empirical evaluation
Deliverable 4	Sustainable infrastructure defense

Pseudo Code:

```
import numpy as np
import time
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import Perceptron
from sklearn.neural_network import MLPClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score
from sklearn.impute import SimpleImputer
np.random.seed(42)
def create_data(n):
    X = np.random.randn(n, 6)
    y = ((X[:, 0] + X[:, 1] - X[:, 2] > 0)).astype(int)
    return X, y
def candidate_elimination(X, y):
    n, d = X.shape
    S = ["0"] * d
    G = [["?"] * d]
    for i in range(n):
        x = X[i]
        if y[i] == 1:
            for j in range(d):
                if S[j] == "0":
                    S[j] = x[j]
                elif S[j] != x[j]:
                    S[j] = "?"
        else:
            new_G = []
            for g in G:
                for j in range(d):
                    if g[j] == "?":
                        h = g.copy()
```

```
                    h[j] = S[j]
                    if h[j] != "?":
                        new_G.append(h)
            G = new_G
    return S, G
X_ce = np.array([
    ["High", "Low", "Normal"],
    ["High", "Low", "High"],
    ["Low", "High", "High"],
    ["High", "Low", "Normal"]
])
y_ce = np.array([1, 1, 0, 1])
S, G = candidate_elimination(X_ce, y_ce)
print("\n--- Candidate Elimination ---")
print("S Boundary:", S)
print("G Boundary:", G)
X, y = create_data(10000)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
tree = DecisionTreeClassifier(criterion="entropy", max_depth=8, random_state=42)
start = time.perf_counter()
tree.fit(X_train, y_train)
tree_pred = tree.predict(X_test)
tree_time = (time.perf_counter() - start) * 1000
print("\n--- Decision Tree ---")
print("Accuracy:", accuracy_score(y_test, tree_pred))
print("Runtime:", round(tree_time, 2), "ms")
perceptron = Perceptron(max_iter=1000, random_state=42)
start = time.perf_counter()
perceptron.fit(X_train, y_train)
p_pred = perceptron.predict(X_test)
p_time = (time.perf_counter() - start) * 1000
print("\n--- Perceptron ---")
print("Accuracy:", accuracy_score(y_test, p_pred))
print("Runtime:", round(p_time, 2), "ms")
```

```
mlp = MLPClassifier(hidden_layer_sizes=(32, 16), activation="relu", max_iter=100, random_state=42)
start = time.perf_counter()
mlp.fit(X_train, y_train)
mlp_pred = mlp.predict(X_test)
mlp_time = (time.perf_counter() - start) * 1000
print("\n--- MLP ---")
print("Accuracy:", accuracy_score(y_test, mlp_pred))
print("Runtime:", round(mlp_time, 2), "ms")
plt.figure()
plt.plot(mlp.loss_curve_)
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.title("MLP Backpropagation Convergence")
plt.grid()
plt.show()
X_missing = X.copy()
mask = np.random.random(X_missing.shape) < 0.05
X_missing[mask] = np.nan
imputer = SimpleImputer(strategy="mean")
X_complete = imputer.fit_transform(X_missing)
nb_em = GaussianNB()
nb_em.fit(X_complete, y)
print("\n--- Missing Data ---")
print("Missing values:", np.isnan(X_missing).sum())
print("Imputed values:", np.isnan(X_complete).sum())
X_train, X_test, y_train, y_test = train_test_split(X_complete, y, test_size=0.2, random_state=42)
nb = GaussianNB()
start = time.perf_counter()
nb.fit(X_train, y_train)
nb_pred = nb.predict(X_test)
nb_time = (time.perf_counter() - start) * 1000
print("\n--- Naive Bayes ---")
```

```
print("Accuracy:", accuracy_score(y_test, nb_pred))
print("Runtime:", round(nb_time, 2), "ms")
def lwr_predict(X_train, y_train, x, tau=1.0):
    distances = np.sum((X_train - x) ** 2, axis=1)
    weights = np.exp(-distances / (2 * tau ** 2))
    return np.sum(weights * y_train) / (np.sum(weights) + 1e-9)
X_lwr = X_train[:500]
y_lwr = y_train[:500]
predictions = []
for x in X_test[:20]:
    predictions.append(lwr_predict(X_lwr, y_lwr, x))
print("\n--- LWR ---")
print("Sample Predictions:")
print(np.round(predictions, 3))
def rbf(x, center, sigma=1.0):
    return np.exp(-np.sum((x - center) ** 2) / (2 * sigma ** 2))
center = X_train.mean(axis=0)
rbf_values = np.array([rbf(x, center) for x in X_train])
print("\n--- RBF ---")
print("Sample RBF values:")
print(rbf_values[:5])
lwr_train = np.array([lwr_predict(X_lwr, y_lwr, x) for x in X_train[:1000]])
X_hybrid = np.column_stack((X_train[:1000], lwr_train))
nb_hybrid = GaussianNB()
nb_hybrid.fit(X_hybrid, y_train[:1000])
lwr_test = np.array([lwr_predict(X_lwr, y_lwr, x) for x in X_test[:100]])
X_hybrid_test = np.column_stack((X_test[:100], lwr_test))
hybrid_pred = nb_hybrid.predict(X_hybrid_test)
print("\n--- Hybrid LWR + Naive Bayes ---")
print("Accuracy:", accuracy_score(y_test[:100], hybrid_pred))
def fitness(k):
    return -(abs(k - 5))
population = np.random.randint(1, 11, 10)
for generation in range(20):
```

```
    population = sorted(population, key=fitness, reverse=True)
    best = population[:5]
    children = []
    for x in best:
        child = x + np.random.choice([-1, 0, 1])
        child = max(1, min(10, child))
        children.append(child)
    population = best + children
best_k = max(population, key=fitness)
print("\n--- Genetic Algorithm ---")
print("Optimized k:", best_k)
def mdl_cost(log_likelihood, parameters, n):
    return -log_likelihood + (parameters / 2) * np.log2(n)
n = len(X_train)
mlp_parameters = sum(w.size for w in mlp.coefs_)
nb_parameters = nb.theta_.size + nb.var_.size + nb.class_prior_.size
mlp_loglikelihood = -mlp.loss_ * n
nb_loglikelihood = -np.mean(-np.log(nb.predict_proba(X_train)[np.arange(len(X_train)),
y_train] + 1e-9)) * n
mlp_mdl = mdl_cost(mlp_loglikelihood, mlp_parameters, n)
nb_mdl = mdl_cost(nb_loglikelihood, nb_parameters, n)
print("\n--- MDL ---")
print("MLP parameters:", mlp_parameters)
print("NB parameters:", nb_parameters)
print("MLP MDL:", mlp_mdl)
print("NB MDL:", nb_mdl)
sizes = [1000, 10000, 100000]
runtime_tree = []
runtime_mlp = []
runtime_nb = []
for n in sizes:
    X, y = create_data(n)
    Xtr, Xte, ytr, yte = train_test_split(X, y, test_size=0.2, random_state=42)
    model = DecisionTreeClassifier(criterion="entropy", max_depth=8)
```

```
    start = time.perf_counter()
    model.fit(Xtr, ytr)
    model.predict(Xte)
    runtime_tree.append((time.perf_counter() - start) * 1000)
    model = MLPClassifier(hidden_layer_sizes=(16,), max_iter=50, random_state=42)
    start = time.perf_counter()
    model.fit(Xtr, ytr)
    model.predict(Xte)
    runtime_mlp.append((time.perf_counter() - start) * 1000)
    model = GaussianNB()
    start = time.perf_counter()
    model.fit(Xtr, ytr)
    model.predict(Xte)
    runtime_nb.append((time.perf_counter() - start) * 1000)
plt.figure()
plt.plot(sizes, runtime_tree, marker="o", label="Decision Tree")
plt.plot(sizes, runtime_mlp, marker="o", label="MLP")
plt.plot(sizes, runtime_nb, marker="o", label="Naive Bayes")
plt.xlabel("Number of Records")
plt.ylabel("Runtime (ms)")
plt.title("Runtime vs Dataset Size")
plt.legend()
plt.grid()
plt.xscale("log")
plt.show()
print("       FINAL SUMMARY")
print("Candidate Elimination: Completed")
print("Decision Tree: Completed")
print("Perceptron: Completed")
print("MLP Backpropagation: Completed")
print("Missing Data Processing: Completed")
print("Naive Bayes: Completed")
print("LWR: Completed")
print("RBF: Completed")
```

print("Hybrid Pipeline: Completed")

print("Genetic Algorithm: Completed")

print("MDL Comparison: Completed")

print("Benchmarking: Completed")

Output:

```
--- Candidate Elimination ---
S Boundary: [np.str_('High'), np.str_('Low'), '?']
G Boundary: [[np.str_('High'), '?', '?'], ['?', np.str_('Low'), '?']]

--- Decision Tree ---
Accuracy: 0.94
Runtime: 32.38 ms

--- Perceptron ---
Accuracy: 0.9955
Runtime: 4.08 ms

Warning (from warnings module):
  File "C:\Users\sekha\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\neural_network\_multilayer_perceptron.py", line 785
    warnings.warn(
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (100) reached and the optimization hasn't converged yet.

--- MLP ---
Accuracy: 0.9965
Runtime: 1195.45 ms
```

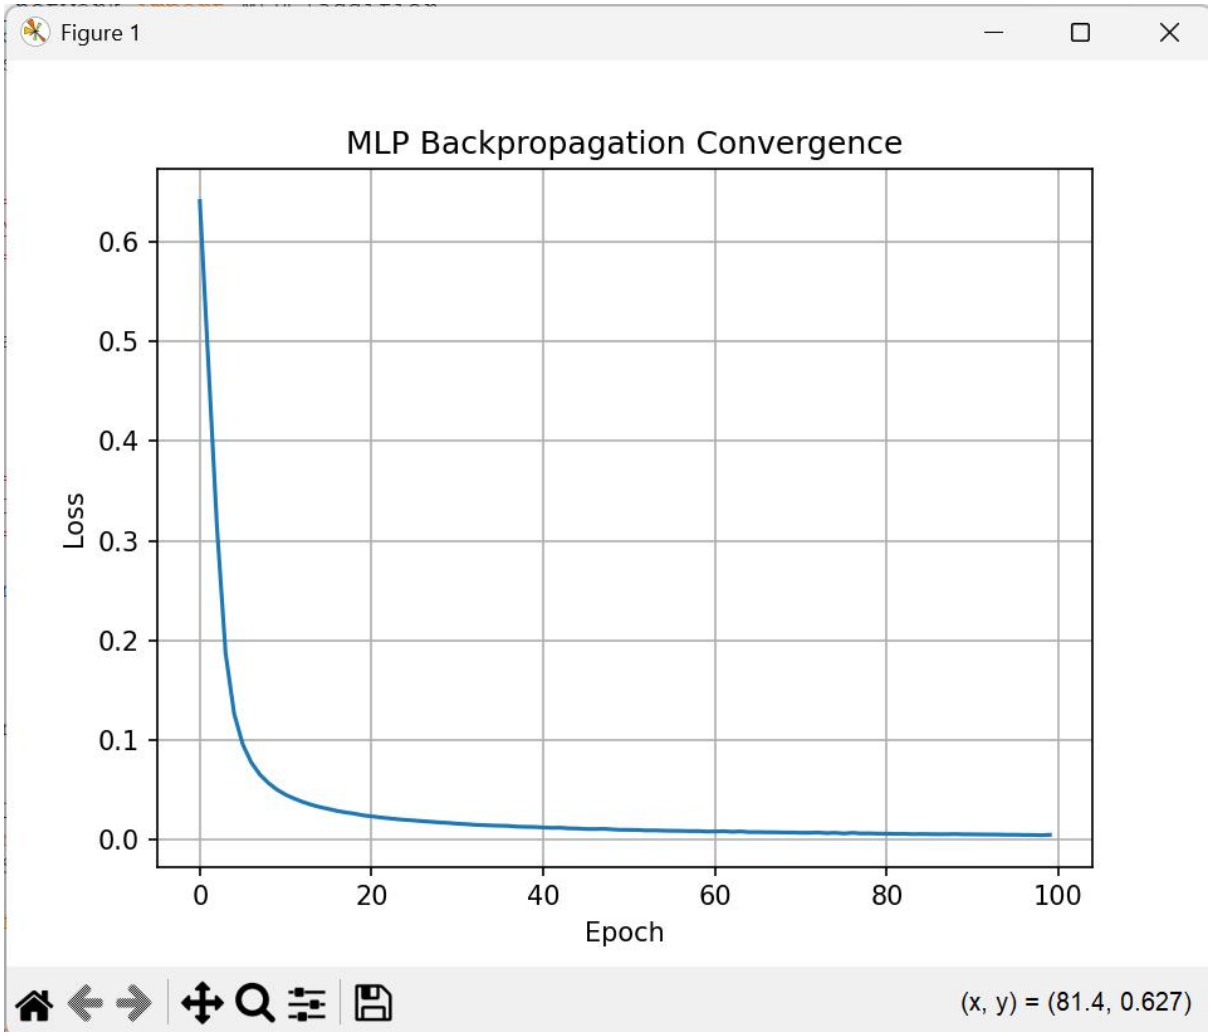

```
--- Missing Data ---
Missing values: 3000
Imputed values: 0

--- Naive Bayes ---
Accuracy: 0.9625
Runtime: 4.77 ms

--- LWR ---
Sample Predictions:
[0.592 0.356 0.085 0.786 0.364 0.274 0.637 0.759 0.495 0.849 0.64  0.627
 0.278 0.632 0.638 0.42  0.378 0.485 0.614 0.535]

--- RBF ---
Sample RBF values:
[0.01131651 0.07140166 0.03372182 0.29707523 0.27362543]

--- Hybrid LWR + Naive Bayes ---
Accuracy: 0.95

--- Genetic Algorithm ---
Optimized k: 5

--- MDL ---
MLP parameters: 720
NB parameters: 26
MLP MDL: 4703.514369386202
NB MDL: 2491.129829170294
```

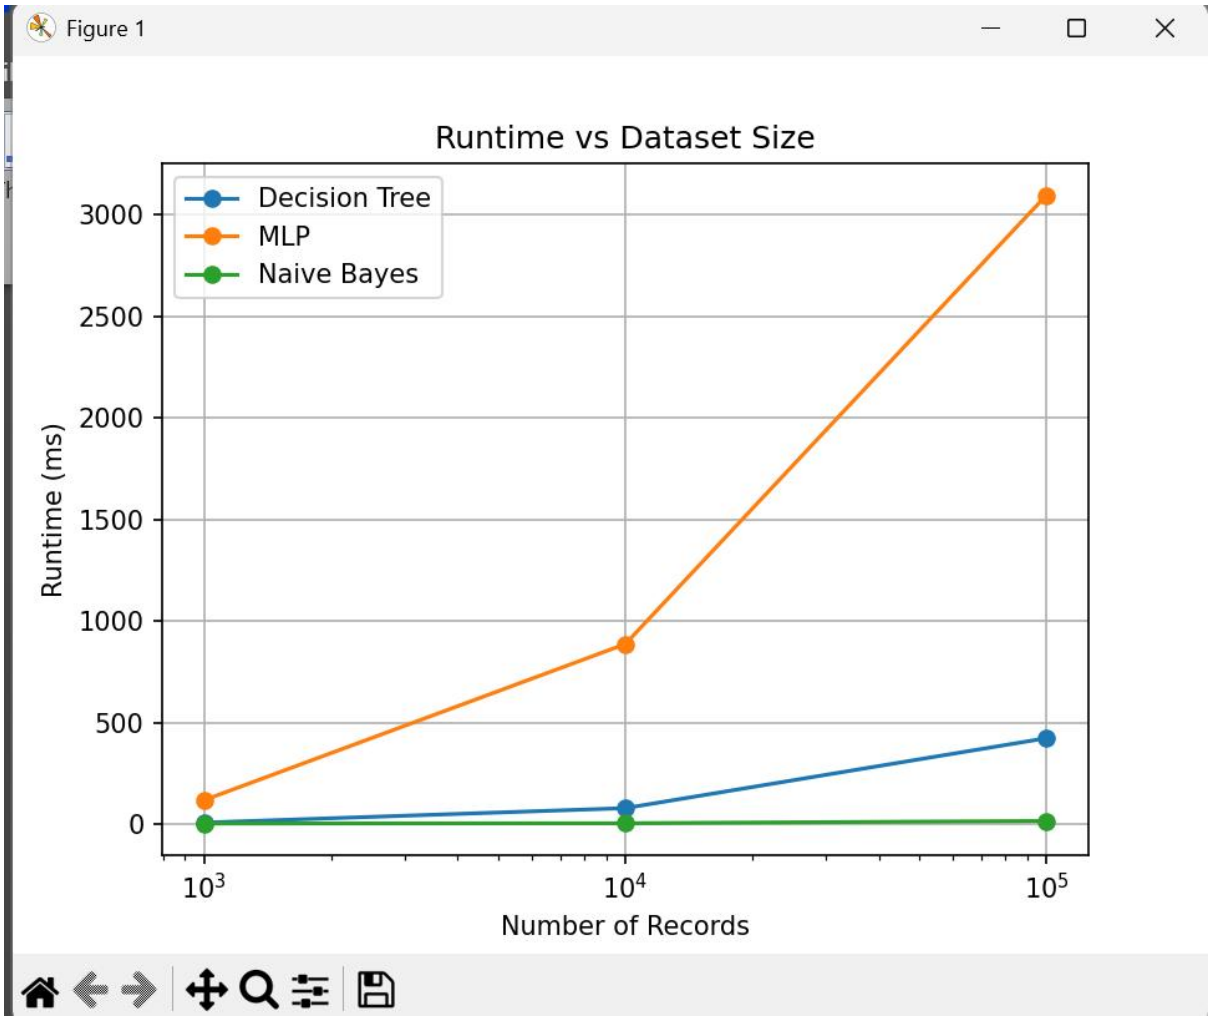

```
================================
       FINAL SUMMARY
================================
Candidate Elimination: Completed
Decision Tree: Completed
Perceptron: Completed
MLP Backpropagation: Completed
Missing Data Processing: Completed
Naive Bayes: Completed
LWR: Completed
RBF: Completed
Hybrid Pipeline: Completed
Genetic Algorithm: Completed
MDL Comparison: Completed
Benchmarking: Completed
================================
>
```